Data, Annotation, and Meaning-Making: The Politics of Categorization in Annotating a Dataset of Faith-based Communal Violence

Mohammad Rashidujjaman Rifat
University of Toronto
Toronto, Ontario, Canada
rifat@cs.toronto.edu

Abdullah Hasan Safir
University of Cambridge
Cambridge, UK
sa2168@cam.ac.uk

Sourav Saha
Shahjalal University of Science and Technology
Sylhet, Bangladesh
sourav95@student.sust.edu

Jahedul Alam Junaed
Shahjalal University of Science and Technology
Sylhet, Bangladesh
jahedul25@student.sust.edu

Maryam Saleki
Fordham University
New York, New York, USA
msaleki@fordham.edu

Mohammad Ruhul Amin
Fordham University
New York, New York, USA
mamin17@fordham.edu

Syed Ishtiaque Ahmed
University of Toronto
Toronto, Ontario, Canada
ishtiaque@cs.toronto.edu

ABSTRACT

Data annotation is a process of meaning-making and is inherently political. The literature on ethics in data-driven technologies explores these political aspects, primarily focusing on questions of bias and power. This paper argues that the politics of annotation often overemphasize secular and modern values and overlooks faith-based, religious, and spiritual aspects (FRS) in data annotation. This oversight particularly affects the postcolonial regions of the Global South, where FRS are intertwined with people's everyday experiences and ethics. We conducted a focus group discussion and contextual inquiries with six annotators who annotated a faith-related "violence" dataset from South Asian YouTube content. Our analysis reveals that FRS blindness in data annotation manifests through the politics of achieving objectivity and the "scientific" process of meaning-making. Due to these goals, which are predominantly shaped by Western values, FRS sensitivities are overlooked from the initial stages of data curation through annotation, ultimately leading to a context collapse within the annotation process. Finally, we advocate for the adaptation of FRS sensitivities into the annotation process and data infrastructure, particularly when the dataset clearly pertains to FRS, to promote greater cultural and contextual inclusivity in annotation practices.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
FAccT '24, June 03–06, 2024, Rio de Janeiro, Brazil
© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 979-8-4007-0450-5/24/06
https://doi.org/10.1145/3630106.3659030

CCS CONCEPTS

• **Human-centered computing → HCI theory, concepts and models**; • **Social and professional topics → Religious orientation**.

KEYWORDS

annotation; meaning-making; decolonizing knowledge practices; communal violence; Faith, Religion, and Spirituality

ACM Reference Format:
Mohammad Rashidujjaman Rifat, Abdullah Hasan Safir, Sourav Saha, Jahedul Alam Junaed, Maryam Saleki, Mohammad Ruhul Amin, and Syed Ishtiaque Ahmed. 2024. Data, Annotation, and Meaning-Making: The Politics of Categorization in Annotating a Dataset of Faith-based Communal Violence. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24), June 03–06, 2024, Rio de Janeiro, Brazil.* ACM, New York, NY, USA, 9 pages. https://doi.org/10.1145/3630106.3659030

1 INTRODUCTION

Data annotation is not merely a technical exercise; it is a process characterized by meaning-making and inherently political in nature [8, 68]. The politics of data annotation have been a focal point of research in AI and data ethics [50, 53, 67, 77]. This body of work has adopted various approaches to study these politics. One strand focuses on issues of justice and harm, particularly concerning biases in data-driven systems [37, 58]. A growing trend of inquiry argues that an exclusive focus on bias may overshadow equally crucial issues related to power dynamics that influence meaning-making in data annotation processes [50]. Researchers in this area have argued that the categorization and interpretation of data are mediated through power-laden infrastructures, where the subjectivity of annotators is often subjugated to imposed, arbitrary categories [50, 51]. These infrastructures are conditioned by a variety of factors, including but not limited to, the pursuit of objectivity,

scientific "neutrality," efficiency, and capitalist logic [68]. Incorporating "power" as a central focus of analysis in these studies has yielded new insights into the interplay between subjectivity and the infrastructures supporting data annotation.

Despite these studies' significant contributions to the politics of annotation and data work, the role of faith, religion, and spirituality (FRS) remains largely unexplored, which is especially important in the data works that explicitly address faith-sensitive issues. This oversight makes the annotation processes predominantly secular practices that neglect FRS-sensitivities. The consideration of FRS is especially important because it often forms a key part of a person's perception of subjectivity [28, 29, 49]. People frequently manifest these subjectivities, often tacitly, when ascribing meaning to their everyday interactions and decision-making processes [20, 59]. As such, acknowledging the role of FRS in human subjectivity and incorporating it into data work is crucial for several reasons. First, FRS aspects are intricately woven and integral to the lives of millions of people, especially in the Global South. These aspects influence both cultural practices and digital interactions [63]. The importance of FRS is evident in scholarly work on related topics. For instance, research on religious discrimination increasingly demonstrates that data-driven methods can be improved by incorporating considerations of faith-related sensitivities [62, 76, 79]. Second, the FRS values of data workers, including annotators, can significantly influence their interpretations and justifications, particularly when handling FRS-sensitive datasets. Since much annotation work is crowdsourced from regions with strong FRS traditions [12], it is especially important to consider how FRS values influence the annotation work. Therefore, integrating FRS perspectives could offer AI ethics studies of data works a more nuanced understanding of the sociocultural factors that affect both annotators and, in turn, the end-users in data-centric systems.

To this end, we conducted a qualitative study to explore the politics and decision-making processes involved in annotating a "violence" dataset. This dataset was created from YouTube comments related to faith-related incidents in South Asian regions. Our research is guided by the question: What knowledge practices, logistical infrastructures, and contextual factors influence the inclusion or exclusion of FRS sensibilities in data collection, cleaning, and annotation processes? Based on focus group discussions and contextual inquiries involving six annotators, we observe that the data annotation politics are heavily influenced by secular and colonial knowledge practices. Key factors shaping this practice include a commitment to "objectivity," established academic norms, expertise hierarchies, and resource scarcity. These processes unfold in a secular context that overlooks FRS considerations, despite the dataset being initially generated through FRS-related keywords. Our findings inspire a critical examination of how data annotation practices can accommodate FRS sensitivities. In summary, the primary contribution of this paper is a critical analysis of data annotation practices within academic settings. Specifically, we focus on how FRS aspects are marginalized as analytical categories within the broader Western and secular politics of data annotation.

2 RELATED WORK

2.1 Politics of Value Alignment in Data Works: Secular and Colonial Practices

A substantial body of literature in ethics in AI employs various modes of inquiry to address the ethical, social, and political aspects in data works. Focusing on questions related to bias, fairness, transparency, power, and identity, a line of inquiry explores the prioritization of values such as profits, standardization, opacity, and corporate interests that shape data annotation practices in the professional and organizational settings of data works [40, 77]. This value alignment often manifests through annotation infrastructures where the erasure of subjectivity and the disregard for annotator diversity are common [39, 77]. A growing line of research pinpoints power dynamics in knowledge practices as the key factor for such marginalization. This body of work provides insights into how organizational hierarchies, the valuing of specialized knowledge by those in power, and business-centric goals contribute to shaping the practices of data annotation [60]. To address the social implications stemming from such practices, existing AI ethics studies often call for greater transparency, the broader inclusion of data workers, improved data education, and the regulatory oversight of data-related activities [6, 33, 47, 75].

While the notion that data annotation is an objective scientific process has been criticized in emerging literature [51] by acknowledging the interplay of various social identities reflected on the sense-making processes involved in annotation work [19, 30, 39, 66], how FRS values are entangled with such processes has largely remained understudied till date. STS researchers have been challenging the notion of a strict separation between science and FRS beliefs for long [9, 22, 42, 65], suggesting they often intermingle but manifest as a separate mode of human existence like politics do [43]. They argue religious and cultural beliefs often shape scientific practices and explore their relevance in subjective interpretations questioning the objectivity of knowledge practices [70]. The Western scientific tradition champions this so-called value neutrality of science, often ignoring diverse worldviews [21] and standpoints [13] of people at the bottom of social hierarchies obstructing their participation in science-making from their marginalized positions.

In recent times, various social media research suggest that FRS values are often reflected on datasets as key characteristics [3, 10], however, we find limited engagement in AI ethics literature that explores the politics of FRS diversities in data annotation. Rather, this body of work is limited in three key ways. First, subjectivity and identity are often framed as obstacles to achieving scientific neutrality and objectivity in data-driven research [14, 44, 57, 78]. Second, recommendations for addressing subjectivity and identity usually conform to Western, scientific, and secular knowledge frameworks, despite the integral role of FRS in shaping people's identities and worldviews [64, 66]. Third, while some research in the Global South does highlight unique contextual challenges, including religious sensitivities [66], it does not further sufficiently explore the impact of FRS considerations on the politics of meaning-making in data work, as evidenced by the lack of further research in the same area.

In response, in this study, we capture FRS sensitivities in the data annotation process and show how AI and data ethics can

benefit from incorporating FRS perspectives in the investigation of annotation politics and data-driven technologies.

2.2 Academic Colonialism, Context Collapse, and Forced Alignment with Foreign Knowledge Practices in Data Works

The AI ethics literature to date primarily reflects perspectives from the Global North [66], while sidelining the ethical diversities and ways of being in the Global South [38, 74]. An absence or scarcity of context-specific insights forces practitioners from the Global South to rely on external knowledge sources, often developed outside their local contexts, to design policies for data collection, preparation, annotation categories, and hiring of annotators, among other aspects of data work. This observation is consistent with what has been termed 'academic colonialism,' wherein knowledge is centralized and flows from specific "virtual centers" to more peripheral regions [1, 2, 34]. These centers establish their influence by generating copious amounts of knowledge, disseminating it through privileged global media channels, and thereby encouraging scholars in peripheral regions to consume this knowledge [4, 69, 71]. This dependency on imported knowledge often persists even when practitioners recognize a contextual mismatch. Moreover, the practitioners are inclined to align with what is often termed the 'circle of academic network' [2], a system generally formulated and controlled by these knowledge centers. Our study will illustrate that similar forms of knowledge dependency are present in data annotation practices in our study site, Bangladesh, a context shaped by postcolonial and postsecular influences. The consequence of this dependency manifests as a "context collapse" [18] in data works. This collapse tends to encourage, as we will demonstrate, a tacit alignment with a homogeneous set of Western and secular values, ethics, and frameworks for meaning-making in annotation.

Contrary to the predominant emphasis on secular values in the discourse on data annotation, various disciplines such as psychology, sociology, and anthropology indicate that FRS elements deeply influence human cognition, affecting everything from ethical decisions to interpretative frameworks [11, 31, 32, 36, 55]. Within AI ethics, an emerging literature examines how FRS intersects with ethical and political considerations in data interpretation [7, 63, 64, 73]. Collectively, these works imply that FRS are not peripheral but foundational to the cognitive frameworks people employ to make sense of the world [26, 59]. Expanding on this, renowned anthropologist Clifford Geertz contends that religion is neither solely a set of supernatural beliefs nor just ritualistic practices [26]. Instead, it offers a framework for interpreting one's experiences, deeply rooted in cultural symbols and shared meanings [26]. Given this substantial evidence of FRS's role in meaning-making, its absence in mainstream, secular discourses on data annotation and AI ethics merits critical scrutiny. The omission of FRS sensitivities not only limits the scope of these discussions but also risks generating a myopic or biased understanding of a dataset, particularly when it comes to annotation practices of FRS-sensitive data.

3 BACKGROUND

We studied the data curation and annotation processes conducted by a group of Bangladeshi researchers [1] aimed at creating an annotated dataset of communal violence in South Asian regions. In the following subsections, we briefly discuss the logistical resources facilitated by the researchers in preparing the dataset, developing the annotation framework, and recruiting the annotators for their research. We should clarify that the study presented in this paper started *after* the researchers had prepared the datasets and already conducted their research.

Dataset preparation: The research group we studied was interested in preparing a communal violence dataset in Bangladesh and West Bengal, India. This regional interest was motivated by their observation that different religious and social groups were using social media and video-sharing platforms to incite violence. The team aimed to build a system for categorizing provocative online content. With this goal, they began by curating a list of significant violent social media incidents from 2012-2022. They surveyed the leading national news outlets as well as utilized their living experiences in the regions to prepare the list. The researchers then searched YouTube videos related to these incidents and finalized a list of 200 pertinent videos based on criteria such as posting date, endorsement of violence, and number of viewers. Finally, utilizing YouTube's public API, they collected over 168,000 comments from these videos, predominantly in the Bangla language. In data preparation, the research team followed standard computational analysis procedures, such as filtering out code-mixed, spammy, and non-Bangla comments, and then anonymizing the data. Employing techniques ranging from unsupervised topic modeling to Guided Latent Dirichlet Allocation (LDA), they narrowed the dataset to approximately 20,000 comments. Lastly, due to time and resource constraints, the researchers selected the first 6,046 comments for human annotation.

Annotation framework development: The next step for the researchers was to design a policy for guiding the data annotation process to categorize the posts. An initial key step in annotation was establishing five categories of violence, four of which were inspired by Galtung's paper on cultural violence written in 1990 [25] and one was proposed by the researchers following a rigorous discussion based on their qualitative understanding of the dataset. The researchers developed a question-based framework to classify the comments into five categories based on their assessed level of associated violence [5]. These levels of violence were rated as "direct violence," "passive violence," or "no violence". This framework finally resulted in a decision tree to better facilitate the annotation process.

Annotator recruitment: Later, the researchers recruited a team of six annotators (all of them were Bangladeshi, four male and two female; three Hindus and three Muslims) to label the dataset. The researchers aimed to form an annotation team with a balanced mix of genders and religious backgrounds. Observing that most collected YouTube comments pertained to Hindu and Islamic religions, a reflection of the demographics in Bangla-speaking regions, they exclusively recruited annotators from these two groups. Of

[1]The term "researchers" specifically denotes the team that conducted data annotation for their study. It is important not to confuse them with the authors of this paper.

particular interest to us were the backgrounds of the annotation team and their annotation works. All six annotators were recruited from a local university's Computer Science (CS) department. The recruitment was carried out by two members of the research team, who were senior undergraduate students at the same department in the university and knew each other before working in the team. Although the annotators worked "independently," they reported substantial agreement on categorizing the comments into one of the five predefined groups. When disagreements involved more than three annotators, they resolved them by consulting an "expert" to make the final decision. The "expert" in this project was one of the project leads in their research team. Initially, the thesis supervisor of these two research team members guided the annotation work. Later, the researchers extended their collaboration beyond their university by partnering with a faculty member at a North American institution. Subsequently, the researchers employed various computational modeling strategies based on linguistic variables and sentence transformers to benchmark their labeled dataset.

We want to clarify that our role as authors of this paper was not to assess the accuracy or efficacy of the dataset, but rather to understand the methodologies employed by the research group and explore how these could be adapted to better reflect local contexts.

4 METHODS

In this paper, we present findings from desk research on the labeled dataset followed by a focus group discussion (FGD) involving the annotator team. We scrutinized the annotated dataset, comprising over 6,046 Bangla YouTube comments, prepared by the researchers as described above. We paid special attention to the comments where annotators had different labels. For a more in-depth analysis, we selected the initial 30 comments and conducted a qualitative examination of their content, aiming to understand the reasons behind annotator disagreements and how they were resolved. We then closely read the research report associated with the dataset, along with the labeling policies prepared by the researchers. The goal of this initial desk research on the materials used in the annotation process and the annotated dataset was to formulate probes for a focus group discussion, which we discuss below.

Following the desk research, we held a three-hour focus group discussion [80] with all six annotators (four male and two female; three Hindus and three Muslims) via Zoom. Conducted in Bangla, the session was audio-recorded with participant consent and later translated and transcribed by our team. We divided the FGD into two parts. In the first part, we asked questions about the annotators' background, their prior experiences of data annotation, their overall experiences of annotating the "violence" dataset, their interpretations of the annotation policies provided by the researchers, their justifications of taking annotation decisions, and their challenges in annotating the dataset. These discussions continued for one hour. For the next two hours of the FGD, we focused on the comments that we selected earlier based on annotators' disagreement and asked specific questions to understand the individual insights and opinions around those comments. We followed up with the annotators later to clarify any ambiguities in their responses from the focus group discussion. Following this focus group, we also conducted unstructured interviews with three researchers in our study. This conversation worked as an opportunity of auto-ethnographic exploration for the researchers. This approach of auto-ethnography has been used as a reflexive research practice in interpretative and qualitative research in which the researchers take the ownership of their subjectivity and its influence in the research methodologies [35, 48]. Due to their auto-ethnographic and reflexive insights, two of the researchers are also co-authors of this paper. The researchers' personal experiences, ranging from data curation to annotation, improved the insights provided by the annotators and helped prevent any misinterpretation of the work conducted by the research group. We were aware of the existing colonial nature of data annotation work [46], and therefore, we did not want our study to be extractive itself. The participants were paid for their participation and time, and the researchers collaborated with us while writing this paper.

For the data analysis phase of our research, we were flexible in using pre-existing theoretical frameworks and focused more on the subtlety and situatedness of the data annotation process. We leveraged the strength of grounded theory [27] by taking inductive approaches to code our qualitative data from the FGD, generating various themes, discussing, iterating the process, and debating the emerging ideas. Such flexible, adaptable, iterative coding techniques and analytical processes helped ensure that our findings are grounded in the data, accurately reflecting the meaning-making experience of a team of data annotators within post-colonial contexts, which are fairly under-studied in literature. We believe our positionalities as authors have played a critical role in the interpretation of the findings reported in this paper. Therefore, we hereby mention that all authors in this study were born and raised in Bangladesh, speak Bangla as their native language, and have been involved in HCI/social and critical data studies research for more than 15 years. This background has helped us capture the nuances of the Bangla dataset, understand Bengali cultural subtleties, and address the data work involved in this research. Finally, after carefully considering the limitations of qualitative methodologies, including researchers' subjectivity, bias, and heuristics, we reached a consensus after several rounds of critical debates within our team to group the findings under the three themes that we discuss next.

5 FINDINGS

Our findings illustrate how the annotation process shapes the categorization of data, notably omitting FRS as values embedded in annotation despite the dataset's dominant focus on faith-based communal violence. This erasure is examined through three key themes that we discuss below.

5.1 Politics of Categories: Secular Knowledge Practices in Devising Annotation Labels

Our study aimed to understand the value judgments of the annotators made during data annotation, focusing primarily on category formation, data labeling, and cleaning. While performing annotation tasks, the project team and annotators often unconsciously and unwillingly navigated the inherent politics involved, resulting in the erasure of "faith" as a category, even when it was strongly visible on the dataset.

The research team we studied analyzed violent comments under YouTube videos from selected news portals in Bangladesh and West Bengal, India. All members, who had computer science backgrounds, followed typical academic norms to adopt a framework-based annotation approach to simplify their work. To develop this framework, the two project leaders (L1 (male) and L2 (male), henceforth) consulted a paper from the field of social and cultural studies [25] as we have mentioned in the background section. While this paper helped the team gain insights into social violence, they were also aware of its limitations. One of the leaders commented:

> *"The paper was written at a time before social media emerged as a phenomenon. For our project we needed to adjust the categories for the present period so that they effectively reflect the threat or violence embedded in YouTube comments by people."* [L2]

Further scrutiny revealed that the project leaders did not consider the paper's contextual differences or its limited understanding of religious conflicts in post-colonial South Asia, except for a brief mention of Mahatma Gandhi's philosophy on cultural violence, also from a Western perspective.

We dug deeper to understand how the project leaders understood the annotation categories. After a few moments of initial consultation between L1 and L2, they reached a consensus on how they interpreted the categories and explained them to us. The team used four categories inspired by the paper: Kill/Attack, Re-socialization/De-socialization/Deportation, Passive Violence/Justification, and Peace/Non-Violence. A fifth category, "Social Rights," emerged from Topic Modelling. They set scores for the categories as follows: Peace/Non-Violence (0), Social Rights (1), Justification/Passive Violence (2), Re-socialization/etc. (3), and Kill/Attack (4). Scores 0-1 denoted non-violence, 2 indicated passive violence, and 3-4 marked direct violence.

The leaders later recruited four annotators (A1, A2, A3 and A4) for labeling the dataset. These annotators only had access to the comments without knowing where they were coming from. It was difficult for them to analyze the texts qualitatively without contexts. Rather, they were given a rigid framework to think quantitatively only with numbers, from 0 to 4. The annotators believed that these numbers made their work easier. For example, one of them said,

> *"The numbering system was quite handy. We were able to assign these five numbers (0 to 4) for all the comments. The issue we often struggled with is the decision of putting what number for a particular comment. Some of them were really confusing."* [A4]

However, despite the team's apparent consensus on annotation categories, our questions about posts such as "Allah eder dhongso koruk" (May Allah destroy them - translated by the authors) revealed varied interpretations among the annotators during the FGD, which is typical in data annotation. The intriguing finding was that the annotators lacked clear guidelines for labeling posts with FRS-sensitivity. Neither L1 nor L2 had consciously engaged with a "faith" category, as they were unsure how to validate it based on the referenced paper or their computational analysis. This may have created a sense of powerlessness among project leaders and annotators, stemming from a lack of rigorous intellectual resources and a reluctance to deviate from the "state-of-the-art" norms in social computing that often dictate academic acceptance or rejection.

While one workaround was to include qualitative methods for the "faith" category to maintain academic rigor, the project leaders emphasized their "scientific" approach for other categories. One team member stated, "We are NLP researchers; we focus more on methods and results than anything else." Our observation is that this focus on mechanical methods and quantifiable results can strip the data of its contextual richness, including FRS-sensitivity.

5.2 Expertise: Faith-blindness in the Inclusion and Exclusion of Experts

Our findings show that the second layer of erasure happens in the inclusion and exclusion politics of expertise. In the project we studied, there was no formal structure for allocating human resources, but a soft hierarchy emerged. Two project leaders (L1 and L2) initially cleaned the dataset and devised the annotation framework, consulting their thesis supervisor (TS). They then recruited four annotators (A1, A2, A3, A4) to score the comments by assessing violence levels without revealing the comments' context. Annotators consulted L1 and L2 for guidance and were compensated financially for their work. The final dataset was shared with an external researcher (PI), under whose supervision L1 and L2 aimed to publish their findings. L1 and L2, as university seniors and subordinates to TS and PI, often hesitate to question or dispute their seniors' directives.

Operating within this soft hierarchy of experts, the annotators were proactive in solving problems that arose during labeling. However, to resolve their confusion, especially regarding the data with religious content, they did not have access to any "expert" or reliable resources. They primarily used readily available online resources. The dataset they were dealing with was rich in religious definitions and cultural references, which were sometimes unclear to the annotators, particularly those from different religious backgrounds. For example, A3 said,

> *"I didn't know some of the Islamic terms that appeared in the comments. I used to google it. I checked the websites that had Quran or Hadith references. Just one or two websites. I think that was enough for our causes."* [A3]

While it is challenging to assess the factual accuracy of the references used, it is worth noting how the annotators conceptualized "accuracy" in their work and how this influenced their judgments. For contextual understanding and cultural interpretation of the texts, which were in Bangla, the annotators did not consult experts like Muslim or Hindu scholars, who could have offered nuanced local contextual knowledge, especially with regard to FRS categories. According to the project leaders, such specialized consultation was deemed unnecessary because the focus was not on the "impact of the violence" [L1] but on determining whether the texts themselves were violent conforming to the guidelines established by the paper they referenced for devising the categories or to their computational analysis. Instead, for fact-checking, the annotators relied on English-language websites that lacked a Bangladesh/South Asian context. Consequently, their perceptions of the inherent violence in the texts, as well as the intensity of such violence, were shaped by these foreign resources.

5.3 Infrastructuring and Drive for Objectivity: Resource Availability, Challenges, and Complicated Perceptions

The third significant factor that marginalized FRS as a category was the team's approach to structuring the data work. Our participants reported making concerted efforts to "avoid biases," according to L1 and L2, and to be as "objective" as possible during the annotation process, as cited by all annotators. However, data labeling is inherently a subjective task requiring individual judgment. To mitigate this, many researchers prioritize ensuring the diversity of annotators, a goal our studied project team also confirmed they aimed to achieve.

We observed that despite best intentions, limited resources—both human and logistical—constrained the team's ability to fully achieve their goals of diversity and objectivity. Originating as a thesis project for L1 and L2, two final-year male undergraduates from different religious backgrounds (Hinduism and Islam), they first aimed for "bias-free" labeling by working separately and resolving disagreements collectively. Initially, they thought their different religions would suffice for creating unbiased annotations. However, they later realized the need for a more systematic approach and recruited four diverse annotators – A1 to A4 – based on religious background, gender, and prior work experience. These annotators, also from the same department and year, were familiar with each other and were chosen largely based on convenience and available resources. Meeting these background criteria and finding appropriate annotators posed a challenge for the project team, especially given that they were university students with relatively limited networks. Consequently, despite the researchers' commitment to diversity, this process led to an annotation team that was ultimately homogeneous in educational and contextual background, with all members lacking expertise in FRS categories.

Moreover, the annotators showed a range of perception of what they meant by "objectivity." We paid closer attention to our annotators while they were sharing their experience of data annotation during our FGD. We found they were speaking in a deeply personal way. For example, one comment they were assessing was *"Apnader moto nari sangbadikra eder theke shikkha niye hijab pore porda kore sangabdikota kora uchit"* (For female journalists like you should learn from them and wear hijab and do journalism - Translated by the authors). One of our participants who assigned level 3 violence (Oppression) according to their agreed framework for this comment said,

> *"It is my right what dress I would wear or not. Who are they to say that?"* [A2]

We highlight the use of "I" in this context because the comment was directed at a female journalist whom the annotators did not personally know. However, the annotator perceived this as a threat to herself, drawing from her own subjective experience as a Muslim woman. Additionally, the annotators mentioned that encountering hate speech targeting their religions adversely affected their emotional and mental well-being during the annotation process. Although the dataset was presented to them without context to minimize bias, the annotators could often deduce the incidents to which the comments referred, especially when those incidents had received significant media coverage. For instance, one participant shared their experience while evaluating a comment expressing hatred toward Hindu communities,

> *"We knew that it was an incident related to finding a Quran in a Pujamandap (Hindu worship place). From the comments I realized how these people think about us (Hindus). These opinions are nothing that didn't exist, I just had direct encounters with these. I felt horrible."* [A1]

However, our participants said that they tried to avoid their "personal biases" during their work and aimed to analyze the texts or comments impartially. This was also reflected in their combined dataset. We found that some texts which were supposed to be offensive for the Hindu community were identified as more violent by the Muslim annotators, possibly with mindsets inclined to overcompensation.

The annotators informed us that they did not disclose their opinions and decisions to each other when they were doing this labeling work so that they would not be influenced by others' opinions. When they needed some clarifications they only reached out to L1 and L2. However, when they independently labeled the intensity of hate speech from 0 to 4, we have found at least one disagreement occurred in 1200+ cases, and many of them had highly contrasting opinions (for example, one annotator assigned 1 and another assigned 4). In these cases, L1 and L2 took the decision based on the majority consensus to produce a "reliable" dataset as they wanted. We also wanted to know how much flexibility the annotators enjoyed while they were working on the dataset. They informed us that they felt they had received adequate time (10 to 12 days) to complete the entire project. They had exams amidst the project and at that time they requested L1 and L2 for more time and they accommodated and adjusted the deadline accordingly. Here, we note how the annotators used their interpersonal relations and communications to negotiate their workload and reduce pressure that could presumably affect their work. From what our participants described during our FGD, we realized that the annotation experiences were like subjective journeys for them, both individually and as a team. However, they were reluctant to admit this; instead, they repeatedly emphasized that their goal was to produce a dataset that would be as "objective" as possible.

Summary of the findings: We depict data annotation as a complex and value-laden process that involves human decision-making at various stages of the work cycle. Although each of the three aspects we have examined is important in its own right, they are interconnected. Collectively, they contribute to the development of a shared sense of meaning among the annotators, which is then conveyed to the end-users of the specific ML/AI project. In this context, we have demonstrated how 'faith' is initially considered in the data collection process but eventually loses its categorical significance due to the overarching goals of objectivity, scientific accuracy, and logistical inconvenience.

6 ADAPTING SENSITIVITY AROUND FAITHS, RELIGIONS, AND SPIRITUALITY IN DATA ANNOTATION

Our findings illustrate critical insights around FRS sensitivity, inclusion and exclusion of related skills and expertise, and tensions

among subjective positionalities within a forced target of objectivity in datawork by the annotators. It captures a wide variety of ways through which the annotators developed a collective sense of meaning-making while categorizing violent posts and how their personal and shared expertise, sentiments, and ethics influenced their work. The nuanced description of data annotation work in relation to FRS-sensitivity presented here challenges the popularly believed assumption that these processes are value-neutral. Additionally, these findings help to reveal how FRS are marginalized as separate categories in data annotation processes within an academic research context. This marginalization primarily occurs through an infrastructural politics of ill-fitting knowledge practices and drives for objectivity and scientism. In this section, we offer critical reflections and recommendations for the AI and data ethics community to include FRS where relevant in data-centric works.

First, we build upon existing literature on data works in the Global South [66, 67] by highlighting the epistemic tension that arises from marginalizing FRS as distinct categories in data annotation. In the context of the "violence" dataset, this tension emerged when FRS were initially considered an explicit category for data collection, only to be later replaced by proxy categories adapted from Galtung's work [25]. We underscore that this replacement was not deliberate but enforced by unspoken power dynamics shaped by academic norms of standardization and objectivity. The implications of this were multifaceted. It deprived annotators of the opportunity to employ FRS as a lens for evaluating the dataset. Indeed, FRS are inevitably intertwined with issues of violence in South Asian countries [17], a context that is familiar to the annotators as well. This omission of FRS led to ambiguity and confusion, particularly when annotators were assessing posts with FRS sentiments. Our analysis showed that YouTube comments containing FRS-related words and phrases were a significant source of annotator disagreement, underlining the need to consider FRS as separate analytical categories in annotation works.

Second, our study demonstrates that the use of foreign and secular knowledge practices in South Asian data work increases the risk of contextual disconnection and potential context collapse. To mitigate this risk, we recommend that the AI ethics community adopt interdisciplinary practices in data work and promote citational justice [41]. To contextualize South Asian research on violence-related data, we recommend engagement with the rich existing literature on violence in the region. This literature addresses the multifaceted aspects of South Asian violence, ranging from the strong influence of FRS to a troubling history of colonialism [15–17, 54]. A survey of this literature could benefit AI ethics research by characterizing violence in South Asia and inspiring further data-centric research. This aligns with the recent decolonial movement in knowledge production [72], wherein the Global South is contributing through innovative methods, theories, and concepts in data-centric research [61].

Finally, we want to clarify that advocating for FRS as distinct analytical categories does not inherently conflict with other categories valuing scientism and objectivity. Instead, we recommend the inclusion of FRS alongside these "modern" categories, aiming for a transition from a universalist framework to a more inclusive pluralistic one in annotation infrastructures. Given that our findings establish both how and why FRS are marginalized, we call for further research to explore mechanisms for integrating FRS into data works. Such integrations will likely lead to ethical tensions, necessitating interdisciplinary approaches that engage with contextual ethics to reconcile conflicts between FRS and other categories in data-centric research.

Taken together, we argue that recognizing and incorporating local and indigenous FRS values in data work should be an important design consideration, which eventually can help ensure that the resulting data-driven systems are more relevant and responsive to the needs, preferences, and cultural contexts of the targeted stakeholders. Thus, our contributions advance HCI scholarship around value sensitive design [23, 24, 56], providing a deeper understanding of the implications of FRS values in data annotation and suggesting how such practices can be more closely aligned with local values, particularly those from marginalized or underrepresented groups in the Global South [45, 52, 64], which may differ from those that are modern, scientific, and Western-centric.

7 CONCLUSION

The goal of this paper is to explore the infrastructural politics that marginalize FRS in annotation work within a dataset focused on South Asian violence. Drawing from our desk research, focus group interviews, and contextual inquiries, we discuss how FRS are marginalized through politics of inclusion, which are dominantly shaped by modern, secular, and objective knowledge practices. We illustrate how non-contextual, yet normatively influential, Western knowledge practices control data work in a South Asian context. We also discuss how resource limitations and a lack of regard for FRS-based expertise lead to FRS blindness in annotation work. Finally, we critically reflect on these findings and recommend a move toward including FRS in data-centric research, thereby initiating a shift toward decolonized and pluralistic data-centric research within AI ethics.

Although we focus on research practices within academic settings in this study, the insights may also provide valuable insights for industry practices. We speculate that industry-academic partnerships can help ensure that data collection and annotation instructions accommodate local contexts, with FRS being one such aspect. Given the emergence of an independent data annotation industry in the Global South [81], we suggest a detailed scrutiny of the recruitment of data workers, the data types they annotate, and their impacts on industry projects. While offering detailed guidelines for integrating FRS considerations would be valuable, the scale and scope of our current study do not fully encompass such comprehensive advice. In our future work, we plan to build upon the critiques presented in this paper and aim to develop policy and design recommendations.

ACKNOWLEDGMENTS

We thank our participants for their valuable insights. Special thanks to our anonymous reviewers, whose constructive comments helped to improve our paper. This research was made possible by the generous grants from the Schwartz Reisman Institute Graduate Fellowship (awarded to Mohammad Rashidujjaman Rifat), the Schwartz Reisman Institute Faculty Fellowship (awarded to Syed Ishtiaque

Ahmed), the Data Sciences Institute Catalyst Grant, the Natural Sciences and Engineering Research Council of Canada (RGPIN-2018-0), the Canada Foundation for Innovation (37608), the Ontario Ministry of Research and Innovation (37608), and Giga Tech, Bangladesh (Enhancement of Bangla Language in ICT - GT/HR/RA/998).

REFERENCES

[1] Syed Farid Alatas. 2000. Academic dependency in the social sciences: Reflections on India and Malaysia. *American Studies International* 38, 2 (2000), 80–96.

[2] Syed Farid Alatas. 2003. Academic dependency and the global division of labour in the social sciences. *Current sociology* 51, 6 (2003), 599–613.

[3] Dina Almanea and Massimo Poesio. 2022. ArMIS-the Arabic misogyny and sexism corpus with annotator subjective disagreements. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference*. 2282–2291.

[4] Claude Alvares. 2011. A critique of Eurocentric social science and the question of alternatives. *Economic and Political Weekly* (2011), 72–81.

[5] Lefteris Jason Anastasopoulos and Jake Ryland Williams. 2019. A scalable machine learning approach for measuring violent and peaceful forms of political protest participation with social media data. *Plos one* 14, 3 (2019), e0212834.

[6] Margherita Antona, George Margetis, Stavroula Ntoa, and Helmut Degen. 2023. Special Issue on AI in HCI. , 1723–1726 pages.

[7] Ananya Bhattacharjee, Dana Kulzhabayeva, Mohi Reza, Harsh Kumar, Eunchae Seong, Xuening Wu, Mohammad Rashidujjaman Rifat, Robert Bowman, Rachel Kornfield, Alex Mariakakis, et al. 2023. Integrating Individual and Social Contexts into Self-Reflection Technologies. In *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–6.

[8] Geoffrey C Bowker and Susan Leigh Star. 2000. *Sorting things out: Classification and its consequences*. MIT press.

[9] John Hedley Brooke. 1991. *Science and religion: Some historical perspectives*. Cambridge University Press.

[10] Pete Burnap and Matthew L Williams. 2016. Us and them: identifying cyber hate on Twitter across multiple protected characteristics. *EPJ Data science* 5 (2016), 1–15.

[11] Heidi A Campbell. 2012. Understanding the relationship between religion online and offline in a networked society. *Journal of the American Academy of Religion* 80, 1 (2012), 64–93.

[12] Mohit Chandra, Manvith Reddy, Shradha Sehgal, Saurabh Gupta, Arun Balaji Buduru, and Ponnurangam Kumaraguru. 2021. " A Virus Has No Religion": Analyzing Islamophobia on Twitter During the COVID-19 Outbreak. In *Proceedings of the 32nd ACM conference on hypertext and social media*. 67–77.

[13] Patricia Hill Collins. 2015. The social construction of black feminist thought. In *Women, knowledge, and reality*. Routledge, 222–248.

[14] Kate Crawford and Trevor Paglen. 2021. Excavating AI: The politics of images in machine learning training sets. *Ai & Society* 36, 4 (2021), 1105–1116.

[15] Veena Das. 2006. *Life and Words: Violence and the Descent into the Ordinary*. Univ of California Press.

[16] Veena Das. 2008. Violence, gender, and subjectivity. *Annual review of anthropology* 37 (2008), 283–299.

[17] Veena Das et al. 1990. Mirrors of violence: Communities, riots, and survivors in South Asia. *(No Title)* (1990).

[18] Jenny L Davis and Nathan Jurgenson. 2014. Context collapse: Theorizing context collusions and collisions. *Information, communication & society* 17, 4 (2014), 476–485.

[19] Mark Díaz, Ian Kivlichan, Rachel Rosen, Dylan Baker, Razvan Amironesei, Vinodkumar Prabhakaran, and Emily Denton. 2022. Crowdworksheets: Accounting for individual and collective identities underlying crowdsourced dataset annotation. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*. 2342–2351.

[20] Avery Dulles. 1984. Faith, Church, and God: Insights from Michael Polanyi. *Theological Studies* 45, 3 (1984), 537–550.

[21] Arturo Escobar. 2018. *Designs for the pluriverse: Radical interdependence, autonomy, and the making of worlds*. Duke University Press.

[22] John H Evans and Michael S Evans. 2008. Religion and science: Beyond the epistemological conflict narrative. *Annu. Rev. Sociol* 34 (2008), 87–105.

[23] Mary Flanagan, Daniel C Howe, and Helen Nissenbaum. 2008. *Embodying values in technology: Theory and practice*. na.

[24] Batya Friedman, Peter H Kahn, Alan Borning, and Alina Huldtgren. 2013. Value sensitive design and information systems. *Early engagement and new technologies: Opening up the laboratory* (2013), 55–95.

[25] Johan Galtung. 1990. Cultural violence. *Journal of peace research* 27, 3 (1990), 291–305.

[26] Clifford Geertz. 2013. Religion as a cultural system. In *Anthropological approaches to the study of religion*. Routledge, 1–46.

[27] Barney G Glaser, Anselm L Strauss, and Elizabeth Strutzel. 1968. The discovery of grounded theory; strategies for qualitative research. *Nursing research* 17, 4 (1968), 364.

[28] Banu Gökarıksel. 2009. Beyond the officially sacred: religion, secularism, and the body in the production of subjectivity. *Social & Cultural Geography* 10, 6 (2009), 657–674.

[29] David J Gouwens. 1996. *Kierkegaard as religious thinker*. Cambridge University Press.

[30] Nitesh Goyal, Ian D Kivlichan, Rachel Rosen, and Lucy Vasserman. 2022. Is your toxicity my toxicity? exploring the impact of rater identity on toxicity annotation. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW2 (2022), 1–28.

[31] Charles Hirschkind. 2001. Civic virtue and religious reason: An Islamic counter-public. *Cultural Anthropology* 16, 1 (2001), 3–34.

[32] Charles Hirschkind. 2006. *The ethical soundscape: Cassette sermons and Islamic counterpublics*. Columbia University Press.

[33] Kori Inkpen, Stevie Chancellor, Munmun De Choudhury, Michael Veale, and Eric PS Baumer. 2019. Where is the human? Bridging the gap between AI and HCI. In *Extended abstracts of the 2019 chi conference on human factors in computing systems*. 1–9.

[34] Lilly Irani, Janet Vertesi, Paul Dourish, Kavita Philip, and Rebecca E Grinter. 2010. Postcolonial computing: a lens on design and development. In *Proceedings of the SIGCHI conference on human factors in computing systems*. 1311–1320.

[35] Alecia Y Jackson and Lisa A Mazzei. 2008. Experience and "I" in autoethnography: A deconstruction. *International review of qualitative research* 1, 3 (2008), 299–318.

[36] Willis Jenkins. 2013. *Ecologies of grace: Environmental ethics and Christian theology*. Oxford University Press.

[37] Kevin Jesse, Premkumar T Devanbu, and Toufique Ahmed. 2021. Learning type annotation: is big data enough?. In *Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. 1483–1486.

[38] Amba Kak. 2020. " The Global South is everywhere, but also always somewhere" National Policy Narratives and AI Justice. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*. 307–312.

[39] Shivani Kapania, Alex S Taylor, and Ding Wang. 2023. A hunt for the Snark: Annotator Diversity in Data Practices. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–15.

[40] Gunay Kazimzade and Milagros Miceli. 2020. Biased priorities, biased outcomes: three recommendations for ethics-oriented data annotation practices. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*. 71–71.

[41] Neha Kumar and Naveena Karusala. 2021. Braving citational justice in human-computer interaction. In *Extended Abstracts of the 2021 CHI Conference on Human Factors in Computing Systems*. 1–9.

[42] Bruno Latour. 2012. *We have never been modern*. Harvard university press.

[43] Bruno Latour. 2013. *An inquiry into modes of existence*. Harvard University Press.

[44] Anita Lavorgna and Pamela Ugwudike. 2021. The datafication revolution in criminal justice: An empirical exploration of frames portraying data-driven technologies for crime prevention and control. *Big Data & Society* 8, 2 (2021), 20539517211049670.

[45] Christopher A Le Dantec, Erika Shehan Poole, and Susan P Wyche. 2009. Values as lived experience: evolving value sensitive design in support of value discovery. In *Proceedings of the SIGCHI conference on human factors in computing systems*. 1141–1150.

[46] Clément Le Ludec, Maxime Cornet, and Antonio A Casilli. 2023. The problem with annotation. Human labour and outsourcing between France and Madagascar. *Big Data & Society* 10, 2 (2023), 20539517231188723.

[47] Duri Long and Brian Magerko. 2020. What is AI literacy? Competencies and design considerations. In *Proceedings of the 2020 CHI conference on human factors in computing systems*. 1–16.

[48] Peter McIlveen. 2008. Autoethnography as a method for reflexive research and practice in vocational psychology. *Australian journal of career development* 17, 2 (2008), 13–20.

[49] Philip A Mellor and Chris Shilling. 2014. Re-conceptualising the religious habitus: Reflexivity and embodied subjectivity in global modernity. *Culture and Religion* 15, 3 (2014), 275–297.

[50] Milagros Miceli, Julian Posada, and Tianling Yang. 2022. Studying up machine learning data: Why talk about bias when we mean power? *Proceedings of the ACM on Human-Computer Interaction* 6, GROUP (2022), 1–14.

[51] Milagros Miceli, Martin Schuessler, and Tianling Yang. 2020. Between subjectivity and imposition: Power dynamics in data annotation for computer vision. *Proceedings of the ACM on Human-Computer Interaction* 4, CSCW2 (2020), 1–25.

[52] Nusrat Jahan Mim. 2021. Gospels of modernity: Digital cattle markets, urban religiosity, and secular computing in the global South. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. 1–17.

[53] Michael Muller, Melanie Feinberg, Timothy George, Steven J Jackson, Bonnie E John, Mary Beth Kery, and Samir Passi. 2019. Human-centered study of data science work practices. In *Extended abstracts of the 2019 CHI conference on human factors in computing systems*. 1–8.

[54] Ashis Nandy. 1988. The politics of secularism and the recovery of religious tolerance. *Alternatives* 13, 2 (1988), 177–194.

[55] Darcia Narvaez, Irene Getz, James R Rest, and Stephen J Thoma. 1999. Individual moral judgment and cultural ideologies. *Developmental psychology* 35, 2 (1999), 478.

[56] Helen Nissenbaum. 2001. How computer systems embody values. *Computer* 34, 3 (2001), 120–119.

[57] Cathy O'neil. 2017. *Weapons of math destruction: How big data increases inequality and threatens democracy*. Crown.

[58] Mihir Parmar, Swaroop Mishra, Mor Geva, and Chitta Baral. 2022. Don't Blame the Annotator: Bias Already Starts in the Annotation Instructions. *arXiv preprint arXiv:2205.00415* (2022).

[59] Michael Polanyi and Harry Prosch. 1975. *Meaning*. University of Chicago Press.

[60] Mohammad Rashidujjaman, Ayesha Bhimdiwala, Ananya Bhattacharjee, Amna Batool, Dipto Das, Nusrat Jahan Mim, Abdullah Hasan Safir, Sharifa Sultana, Taslima Akter, C Estelle Smith, et al. 2023. Many Worlds of Ethics: Ethical Pluralism in CSCW. (2023).

[61] Noopur Raval. 2019. An agenda for decolonizing data science. *Spheres: Journal for Digital Cultures* 5 (2019), 1–6.

[62] Mohammad Rashidujjaman Rifat, Ashratuz Zavin Asha, Shivesh Jadon, Xinyi Yan, Shion Guha, and Syed Ishtiaque Ahmed. 2024. Combating Islamophobia: Compromise, Community, and Harmony in Mitigating Harmful Online Content. *ACM Transactions on Social Computing* (2024).

[63] Mohammad Rashidujjaman Rifat, Firaz Ahmed Peer, Hawra Rabaan, Nusrat Jahan Mim, Maryam Mustafa, Kentaro Toyama, Robert B Markum, Elizabeth Buie, Jessica Hammer, Sharifa Sultana, et al. 2022. Integrating Religion, Faith, and Spirituality in HCI. In *CHI Conference on Human Factors in Computing Systems Extended Abstracts*. 1–6.

[64] Mohammad Rashidujjaman Rifat, Toha Toriq, and Syed Ishtiaque Ahmed. 2020. Religion and sustainability: Lessons of sustainable computing from Islamic religious communities. *Proceedings of the ACM on Human-Computer Interaction* 4, CSCW2 (2020), 1–32.

[65] Holmes Rolston. 2006. *Science and religion: A critical survey*. Number 188. Templeton Foundation Press.

[66] Nithya Sambasivan, Erin Arnesen, Ben Hutchinson, Tulsee Doshi, and Vinodkumar Prabhakaran. 2021. Re-imagining algorithmic fairness in india and beyond. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 315–328.

[67] Nithya Sambasivan, Shivani Kapania, Hannah Highfill, Diana Akrong, Praveen Paritosh, and Lora M Aroyo. 2021. "Everyone wants to do the model work, not the data work": Data Cascades in High-Stakes AI. In *proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. 1–15.

[68] Morgan Klaus Scheuerman, Alex Hanna, and Emily Denton. 2021. Do datasets have politics? Disciplinary values in computer vision dataset development. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2 (2021), 1–37.

[69] Caroline M Schöpf. 2020. The Coloniality of Global Knowledge Production: Theorizing the Mechanisms of Academic Dependency. *Social Transformations: Journal of the Global South* 8, 2 (2020).

[70] Helaine Selin. 2013. *Encyclopaedia of the history of science, technology, and medicine in non-westen cultures*. Springer Science & Business Media.

[71] Cheng Feng Shih. 2010. Academic colonialism and the struggle for indigenous knowledge systems in Taiwan. *Social Alternatives* 29, 1 (2010), 44–47.

[72] Linda Tuhiwai Smith. 2021. *Decolonizing methodologies: Research and indigenous peoples*. Bloomsbury Publishing.

[73] Sharifa Sultana and Syed Ishtiaque Ahmed. 2019. Witchcraft and hci: Morality, modernity, and postcolonial computing in rural bangladesh. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*. 1–15.

[74] Scott Timcke. 2023. Towards a More Comprehensive AI Ethics: How Global South Perspectives Can Enrich Current Approaches to AI Governance. *Scott Timcke and Theresa Schültken (2023) Towards a More Comprehensive AI Ethics: How Global South Perspectives Can Enrich Current Approaches to AI Governance, UN Unite Brief* (2023).

[75] Tânia Valente. 2023. Future Visions for a Decolonized Future of HCI. In *HCI International 2023 Posters: 25th International Conference on Human-Computer Interaction, HCII 2023, Copenhagen, Denmark, July 23–28, 2023, Proceedings, Part I*. Springer Nature, 109.

[76] Bertie Vidgen and Taha Yasseri. 2020. Detecting weak and strong Islamophobic hate speech on social media. *Journal of Information Technology & Politics* 17, 1 (2020), 66–78.

[77] Ding Wang, Shantanu Prabhat, and Nithya Sambasivan. 2022. Whose AI Dream? In search of the aspiration in data annotation.. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems*. 1–16.

[78] Jutta Weber and Bianca Prietl. 2021. AI in the age of technoscience: On the rise of data-driven AI and its epistem-ontological foundations. In *The Routledge social science handbook of AI*. Routledge, 58–73.

[79] Susann Wiedlitzka, Gabriele Prati, Rupert Brown, Josh Smith, and Mark A Walters. 2023. Hate in word and deed: the temporal association between online and offline islamophobia. *Journal of quantitative criminology* 39, 1 (2023), 75–96.

[80] Sue Wilkinson. 1998. Focus group methodology: a review. *International journal of social research methodology* 1, 3 (1998), 181–203.

[81] Angie Zhang, Alexander Boltz, Chun Wei Wang, and Min Kyung Lee. 2022. Algorithmic management reimagined for workers and by workers: Centering worker well-being in gig work. In *Proceedings of the 2022 CHI conference on human factors in computing systems*. 1–20.